

INTRODUCTION TO PARALLEL PROGRAMMING

Caspar van Leewen
Carlos Teijeiro Barjas

Goals

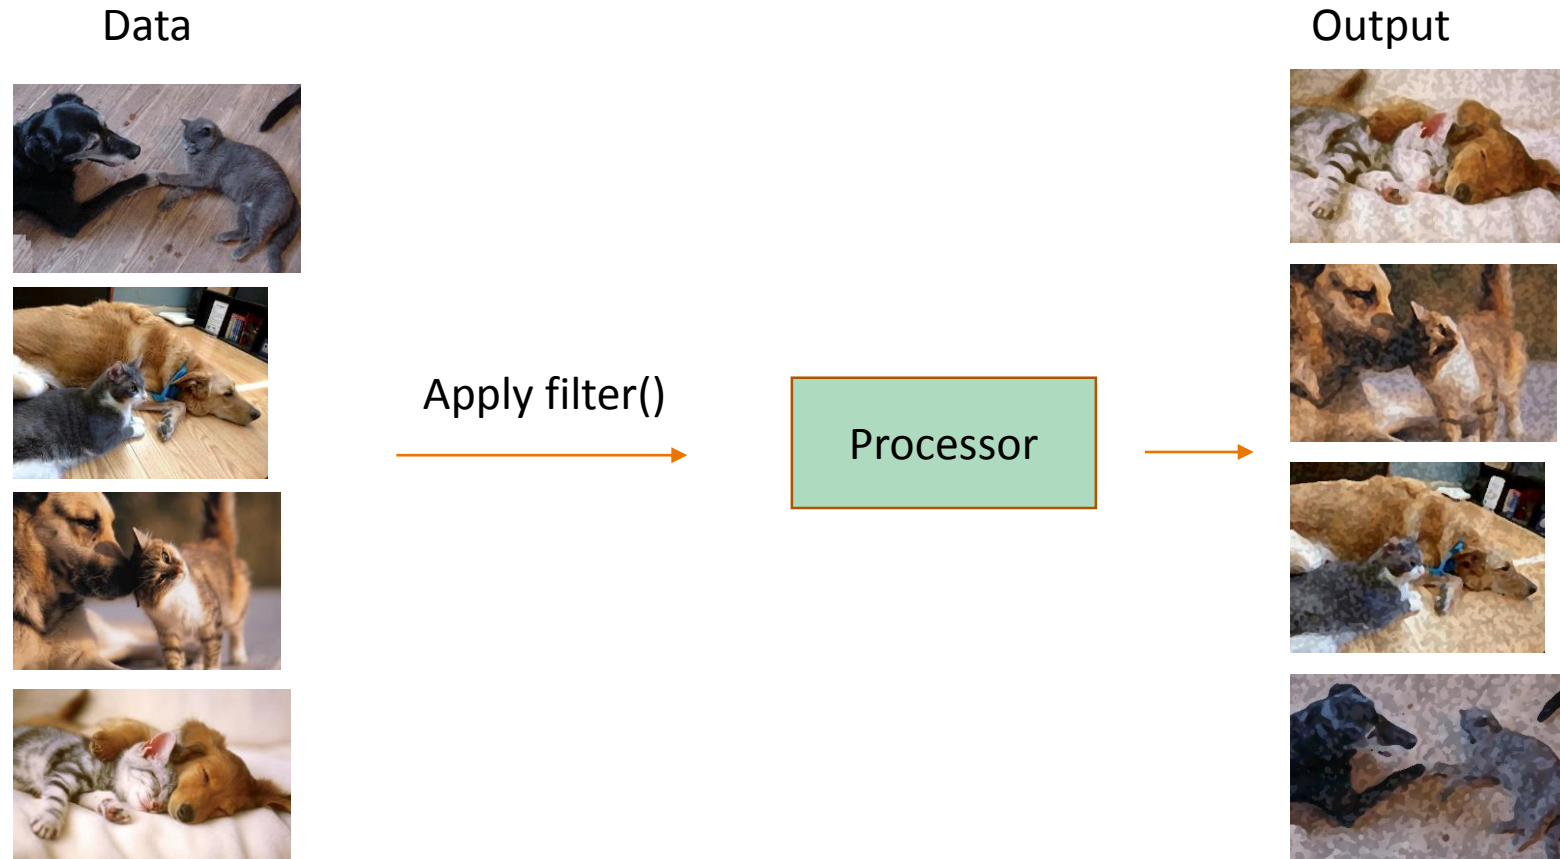
- The goal is to understand:
 - Merits and limits of parallel computing
 - Parallel programming models (task / data parallelism)
 - Differences between shared and distributed memory systems

Content

- What is parallel computing?
- Parallel programming models
- Different types of systems

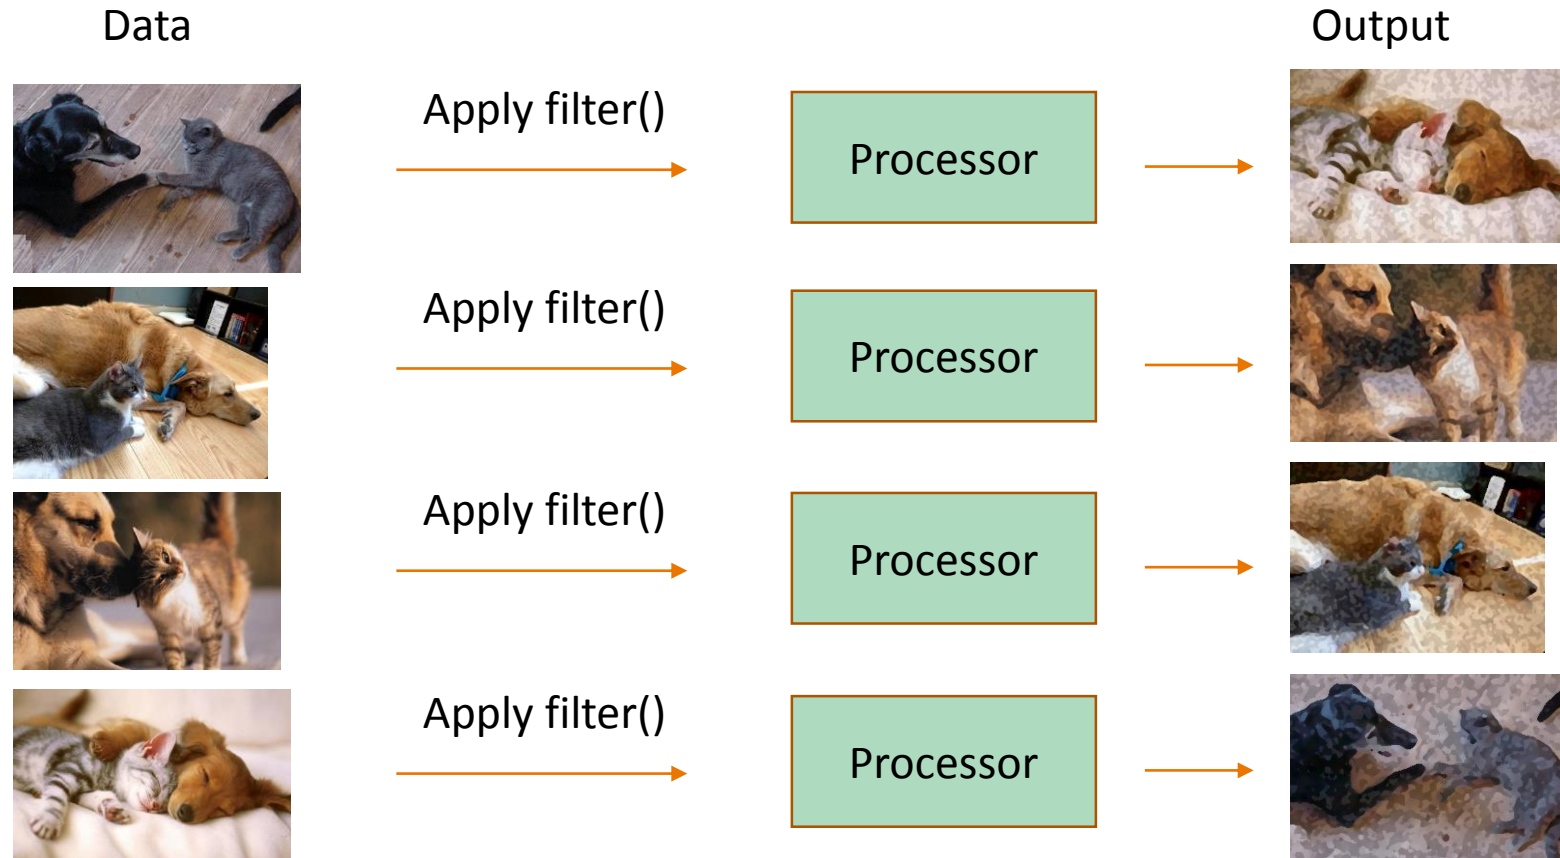
What is parallel computing?

- Parallel computing
 - Multiple processors or computers working on a single computational problem



What is parallel computing?

- Parallel computing
 - Multiple processors or computers working on a single computational problem



Parallel computing: benefits

- Benefits
 - Solve computationally intensive problems (speedup)
 - Solve problems that don't fit a single memory (multiple computers)
- Requirements:
 - Problem should be divisible in smaller tasks

Real problem: often partly parallelizable

Serial program

Non-parallelizable
(e.g. initialization)

Parallelizable computation (e.g. filters)

Non-parallelizable

Parallel program (e.g. 2 processors)

Non-parallelizable
(e.g. initialization)


Parallelizable
computation (e.g.
filters)

Non-parallelizable


Time

Real problem: often partly parallelizable

Serial program



Parallel program (e.g. 50 processors)




Amdahl's Law

$$Speedup = \frac{1}{(1-p) + \frac{p}{N}}$$

N = number of processors

p = parallel fraction



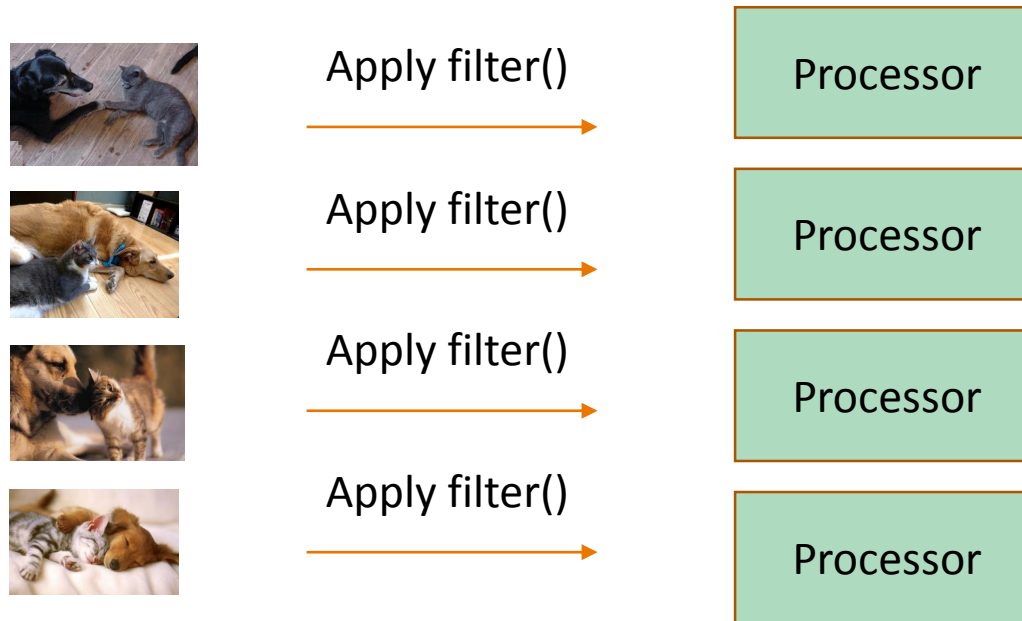
Amdahl's law describes *strong scaling*

- Strong scaling
 - Variation of solution time with #processors for fixed *total* problem size
 - Possibilities to run the same problem in shorter time



Amdahl's law describes *strong scaling*

- Strong scaling
 - Variation of solution time with #processors for fixed *total* problem size
 - Possibility to run the same problem in shorter time



Amdahl's law is theoretical...

- ... and reality is often worse!
- Further issues to tackle
 - Load balancing
 - Process scheduling
 - Communication overhead
 - File Input/Output
 - ...

Do we have hope...???

Gustafson's Law

$$\text{Speedup} = 1 + p \cdot (N - 1)$$

N = number of processors

p = parallel fraction




Image source: https://computing.llnl.gov/tutorials/parallel_comp

Gustafson's law describes *weak scaling*

- Weak scaling
 - Variation of solution time with #processors for fixed problem size *per processor*
 - Possibilities to run a bigger problem in the same time



Gustafson's law describes *weak scaling*

- Weak scaling
 - Variation of solution time with #processors for fixed problem size *per processor*
 - Possibilities to run a bigger problem in the same time




Strong scaling vs. weak scaling

- Strong scaling: fixed work

$$Speedup = \frac{1}{(1 - p) + \frac{p}{N}}$$


p=0.5



- Weak scaling: fixed work per processor

$$Speedup = 1 + p \cdot (N - 1)$$

p=0.5



Weak scaling is often the most relevant to HPC computations

- Some examples
 - Physics: "I can only run my fluid simulation a small domain / low resolution on my local PC"
 - Chemistry: "I can simulate a small molecule on my PC, but I want to simulate a big one"
- Common background
 - A big system / molecule increases the *total* work
 - Distributing larger work over multiple processors keeps the work *per processor* constant
 - ... this is weak scaling!

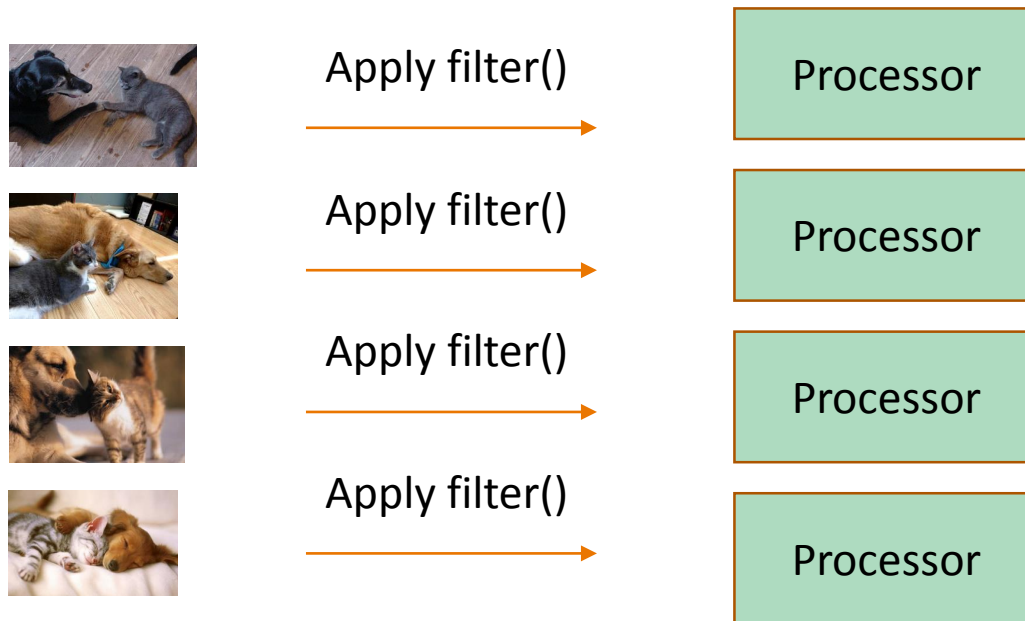
So now how can we program in parallel?

Parallel programming models

- “How do I parallelize my code?”
- Two well-known programming models
 - Data parallelism
 - Task parallelism

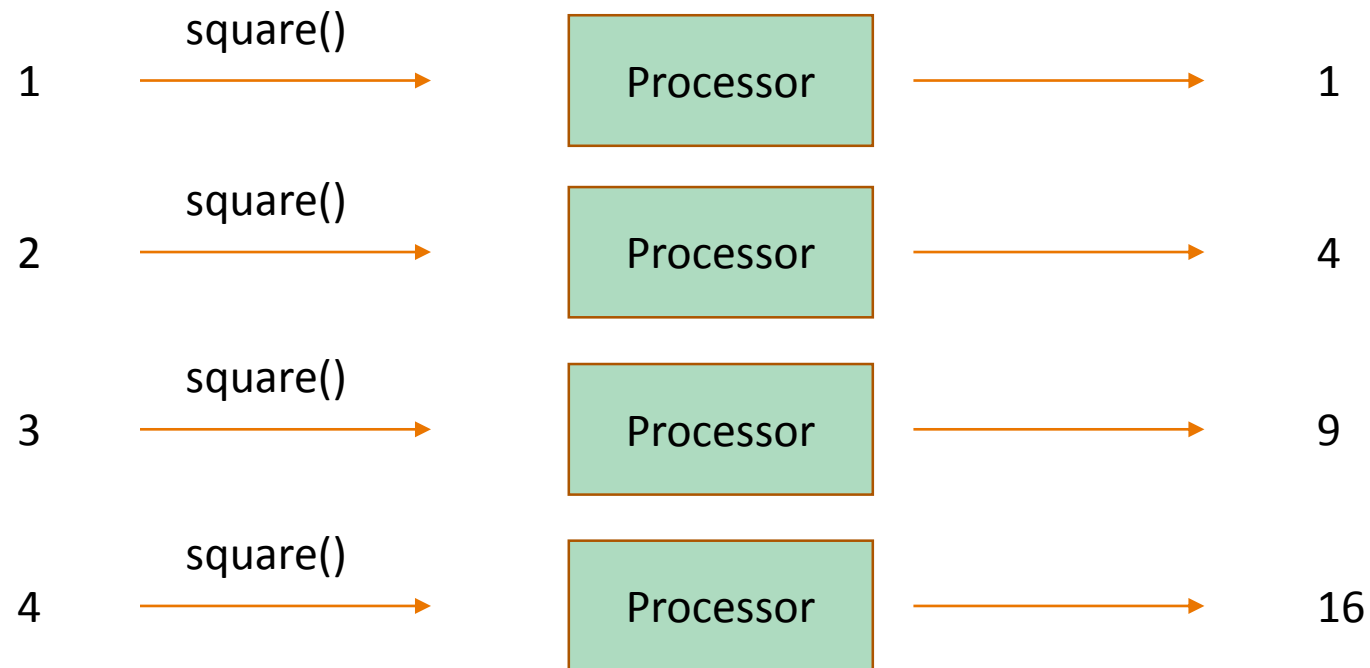
Parallel programming models: data parallelism

- Each processor performs the same task on different data



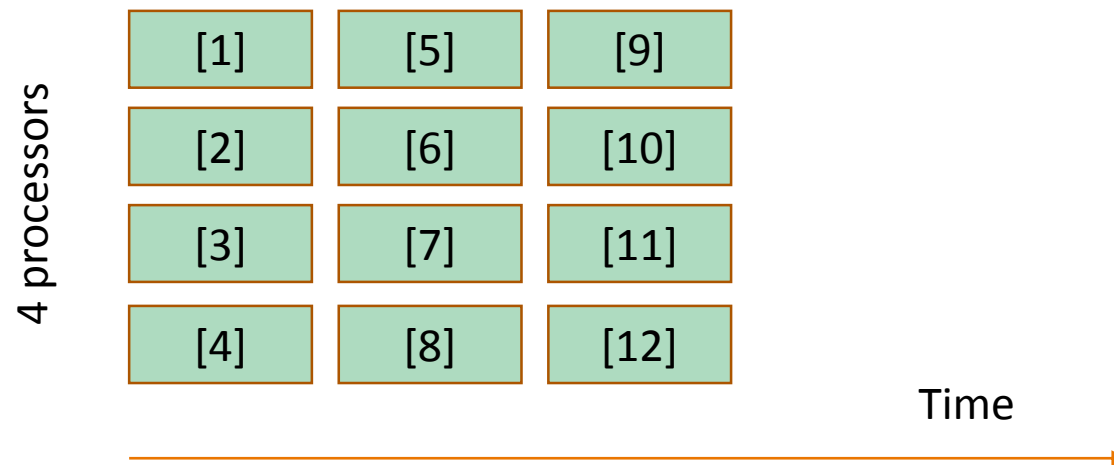
Parallel programming models: data parallelism

- Each processor performs the same task on different data



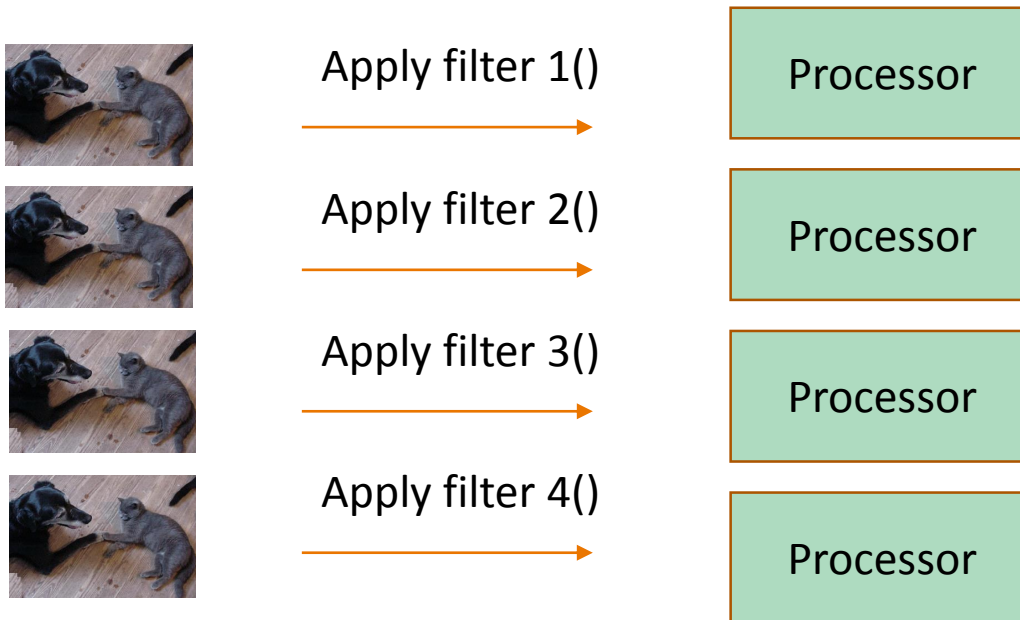
Parallel programming models: data parallelism

- Amount of parallelization depends on input data size
- Load balancing may be relatively easy
 - Same task on each data element
 - Approximately same time per element



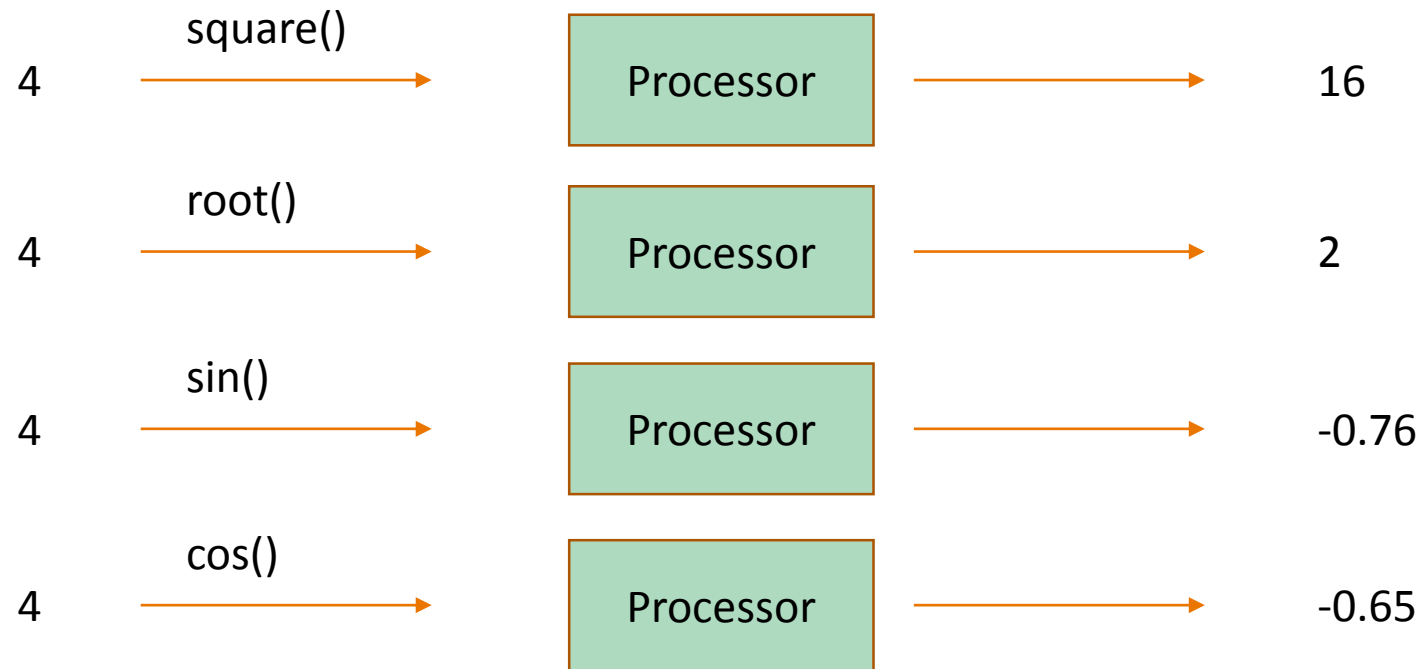
Parallel programming models: task parallelism

- Each processor performs a different task on the same data



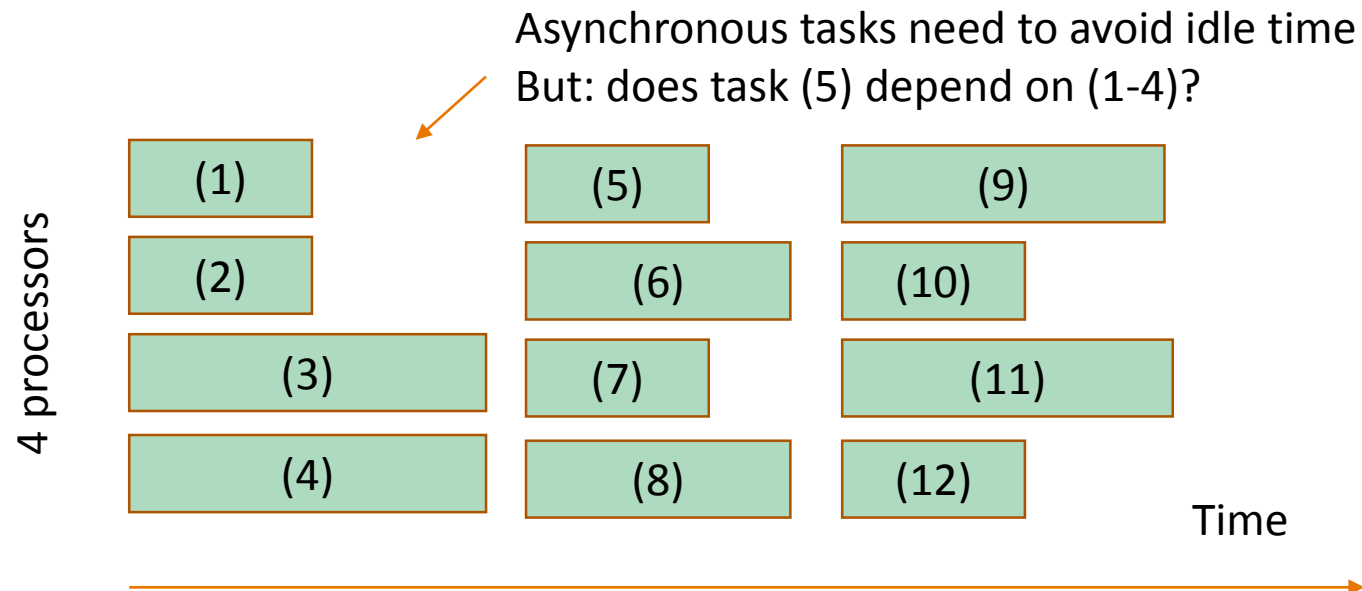
Parallel programming models: task parallelism

- Each processor performs a different task on the same data



Parallel programming models: task parallelism

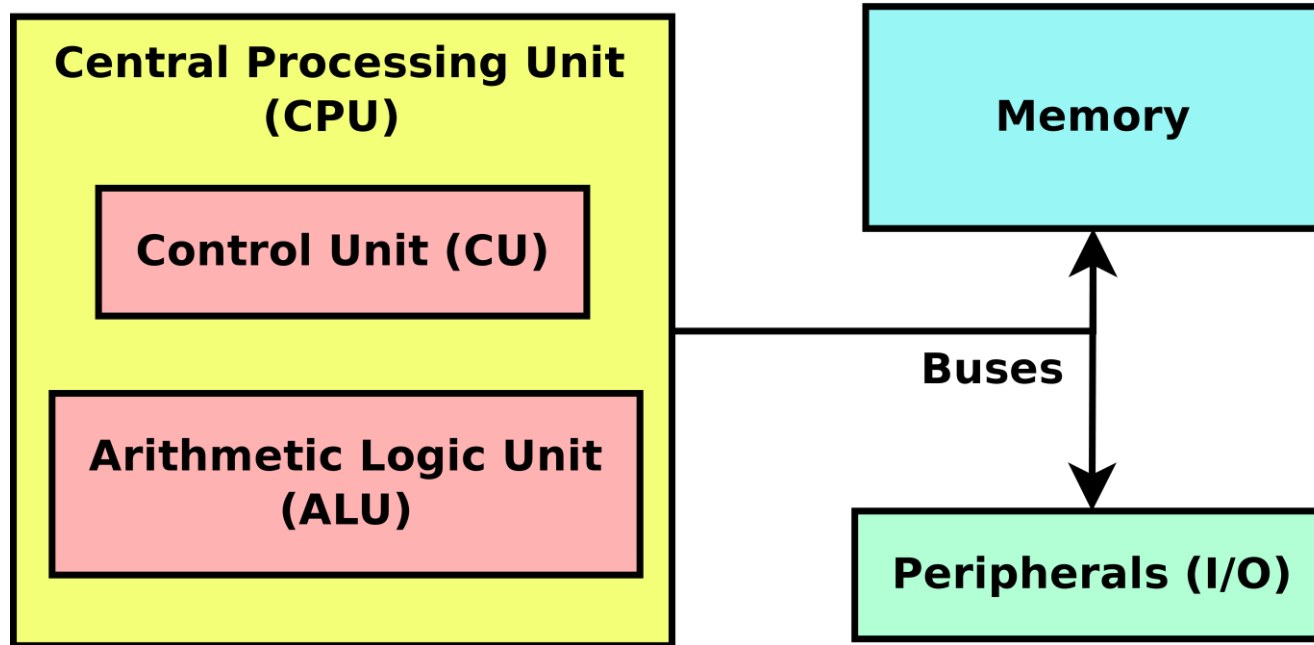
- Amount of parallelization depends on the number of tasks
- Load balancing can be very difficult
 - Heterogeneous tasks may be executed over the same data
 - Each task may take a very different amount of time



Theory is essential to understand concepts...

**but parallel programming also requires
understanding of the systems!**

A computer is...



A memory is...


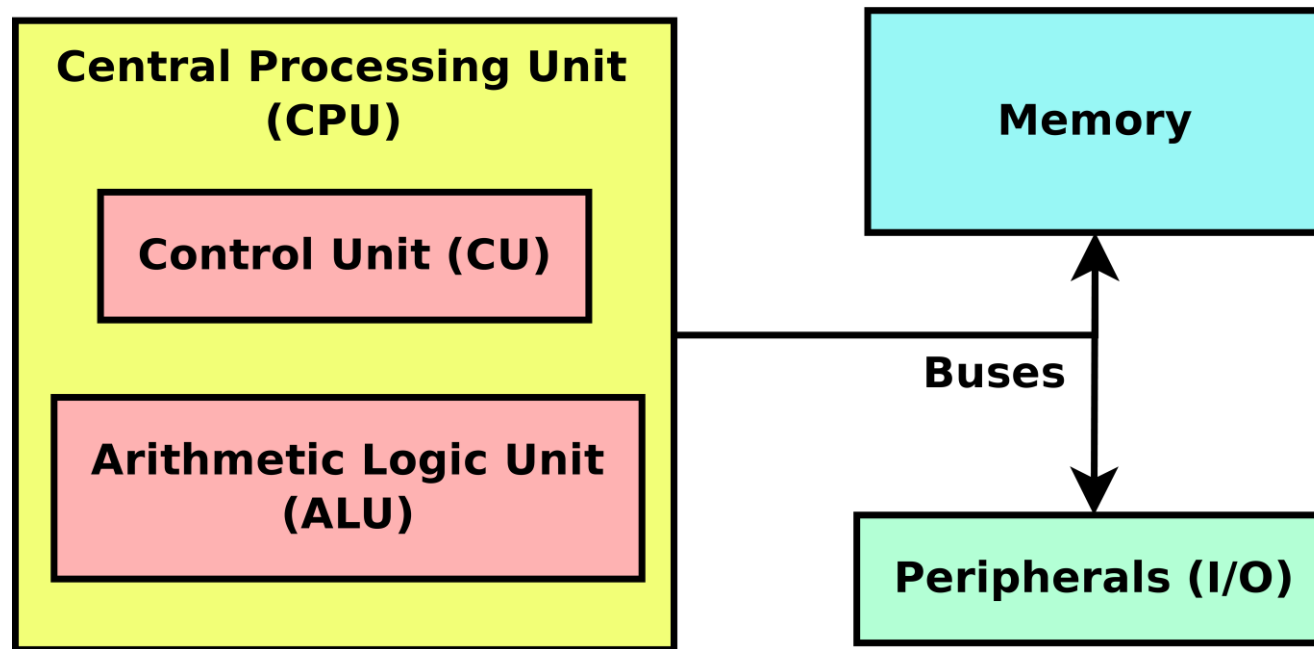


Image source: <https://upload.wikimedia.org/wikipedia/commons/c/c0/8 bytes vs. 8Gbytes.jpg>

A computer is...



A CPU is...

Example of the MIPS architecture: control unit lines in orange, data lines in black

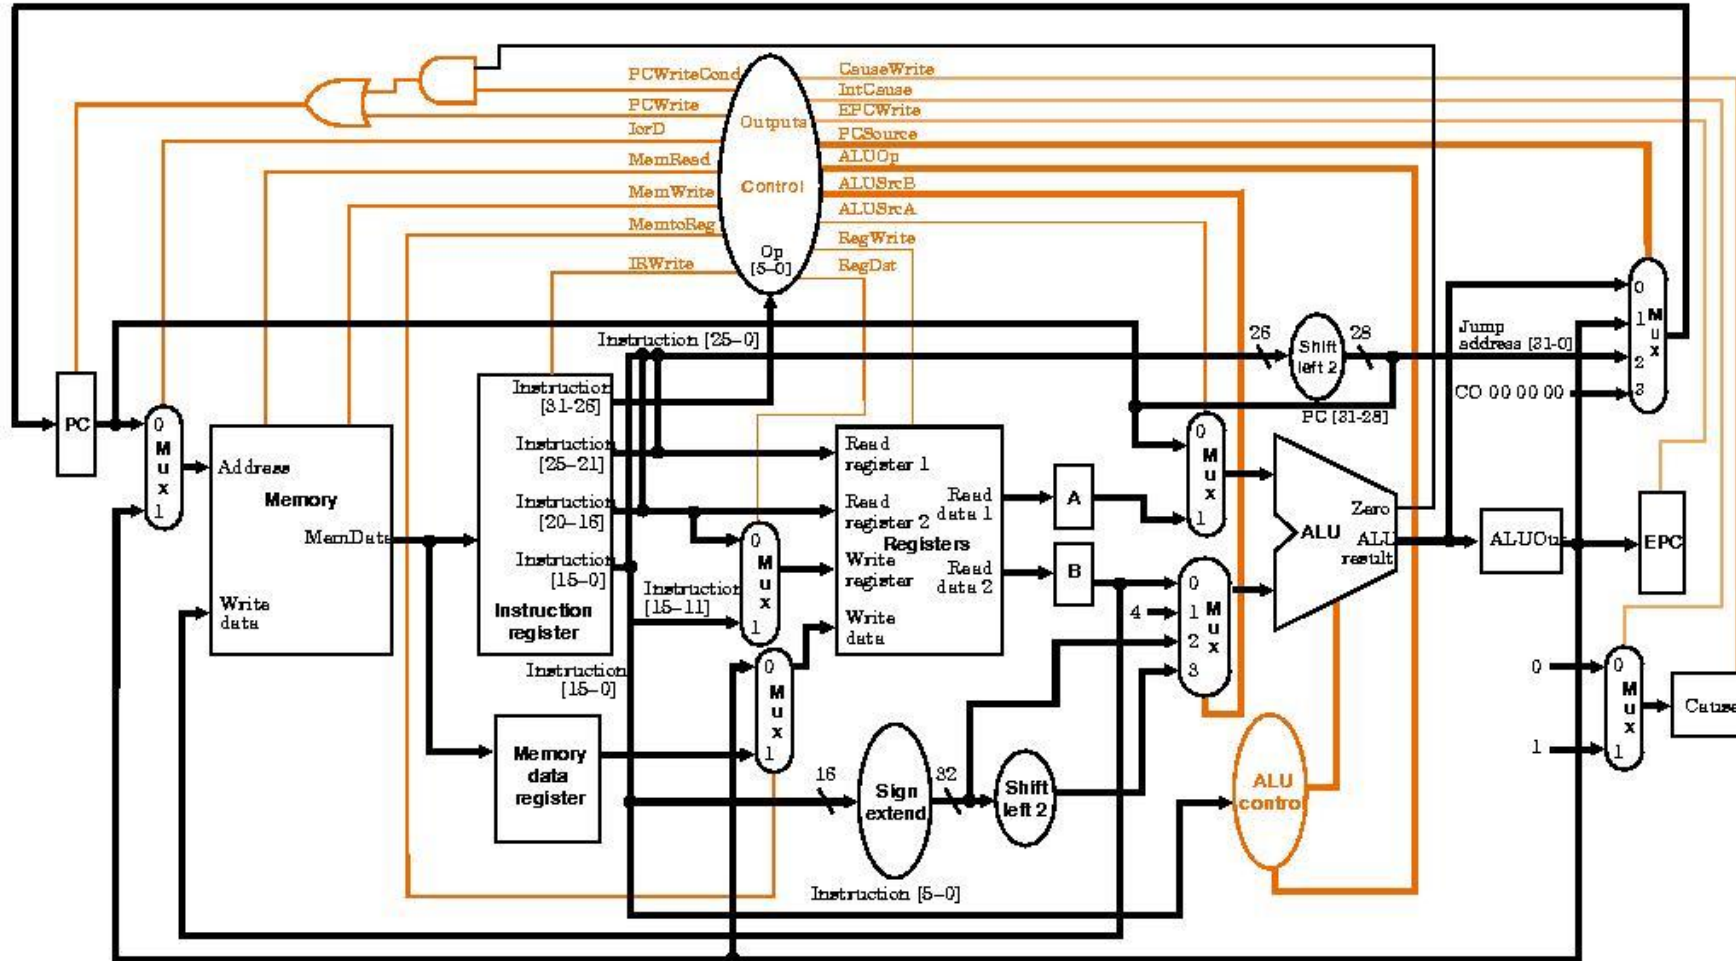


Image source: <http://people.cs.pitt.edu/~don/coe1502/current/Unit4a/Unit4a.html> (last visited: 2018)

A program becomes a process...

SOURCE CODE

```
void main() {  
    int id = 0;  
    printf("hello world from %d !\n",id);  
}
```

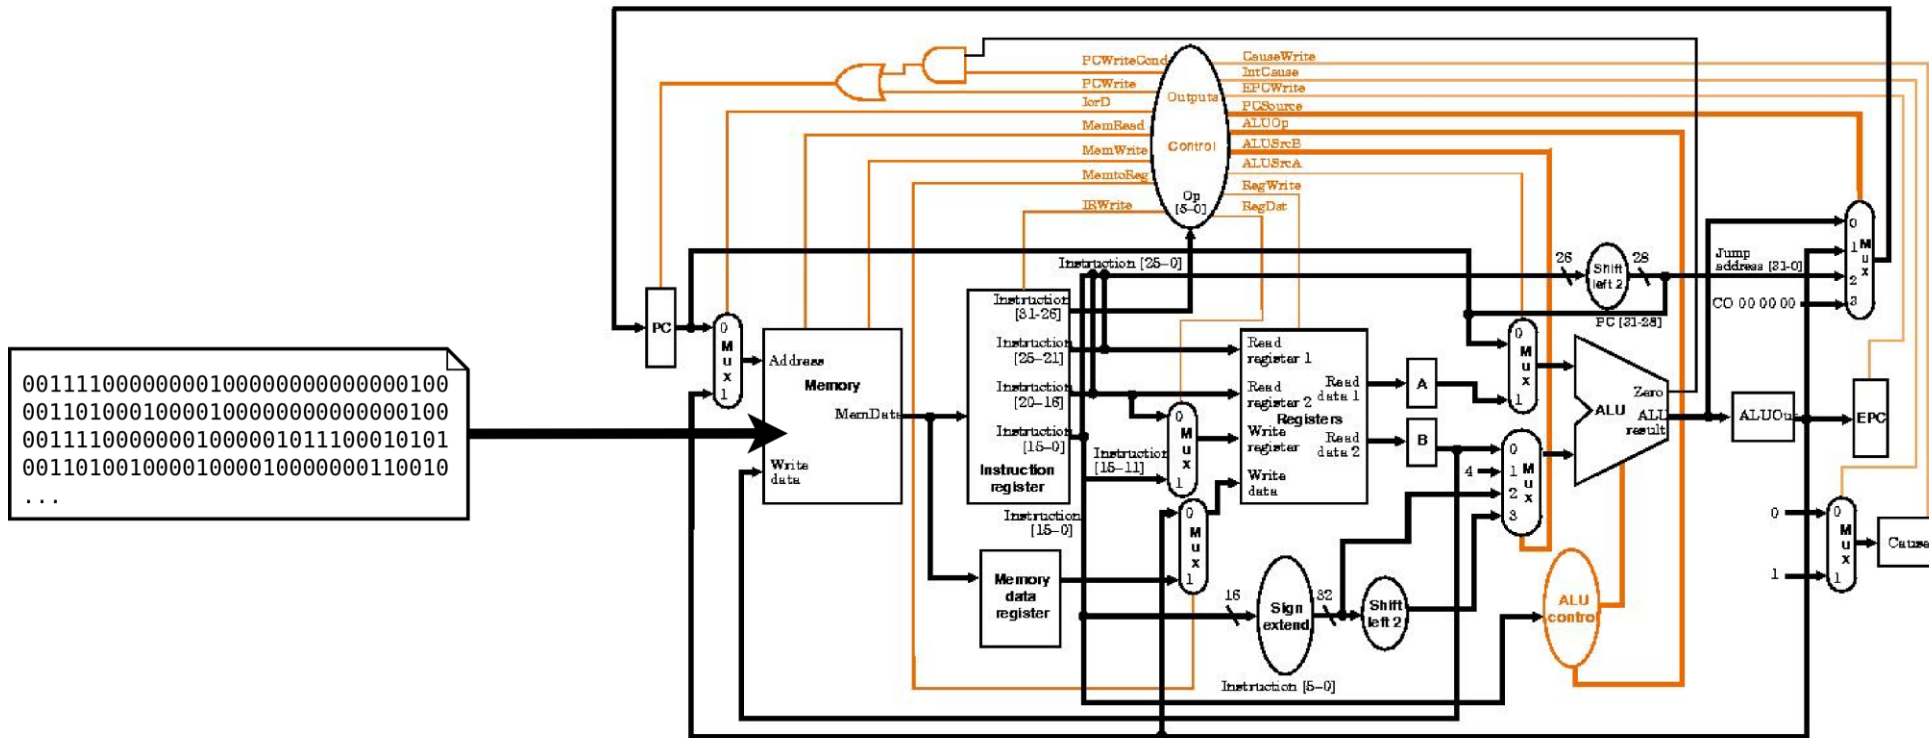
ASSEMBLY CODE

```
.data  
str: .asciiz "hello world from "  
...  
  
.text  
main:  
    li $v0, 4  
    la $a0, str  
    ...  
    syscall  
    ...
```


```
00111100000000100000000000000100  
00110100010000100000000000000100  
00111100000001000001011100010101  
00110100100001000010000000110010  
...
```

MACHINE CODE

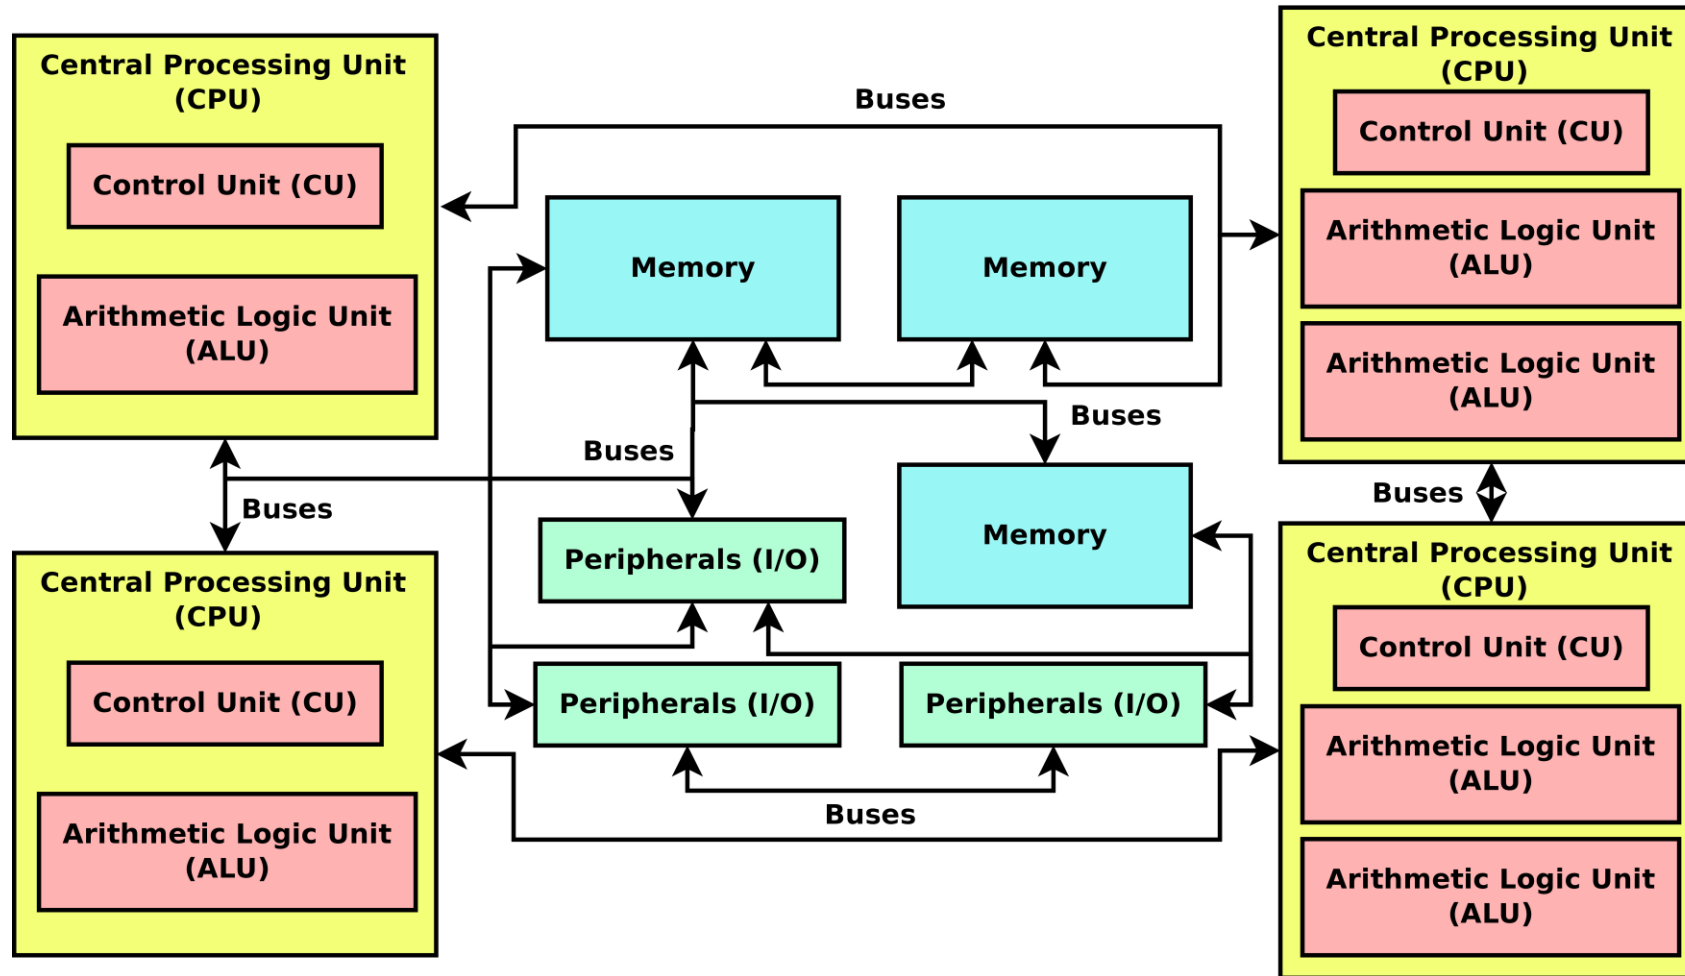
A program becomes a process...



Different types of computer architectures




Different types of computers architectures (!!1!!!!11!!eleven!!)




Different types of computer architectures

- Shared memory: all processors access the same memory (e.g. typical laptop)
 - Early model: uniform memory access for every processor (UMA)



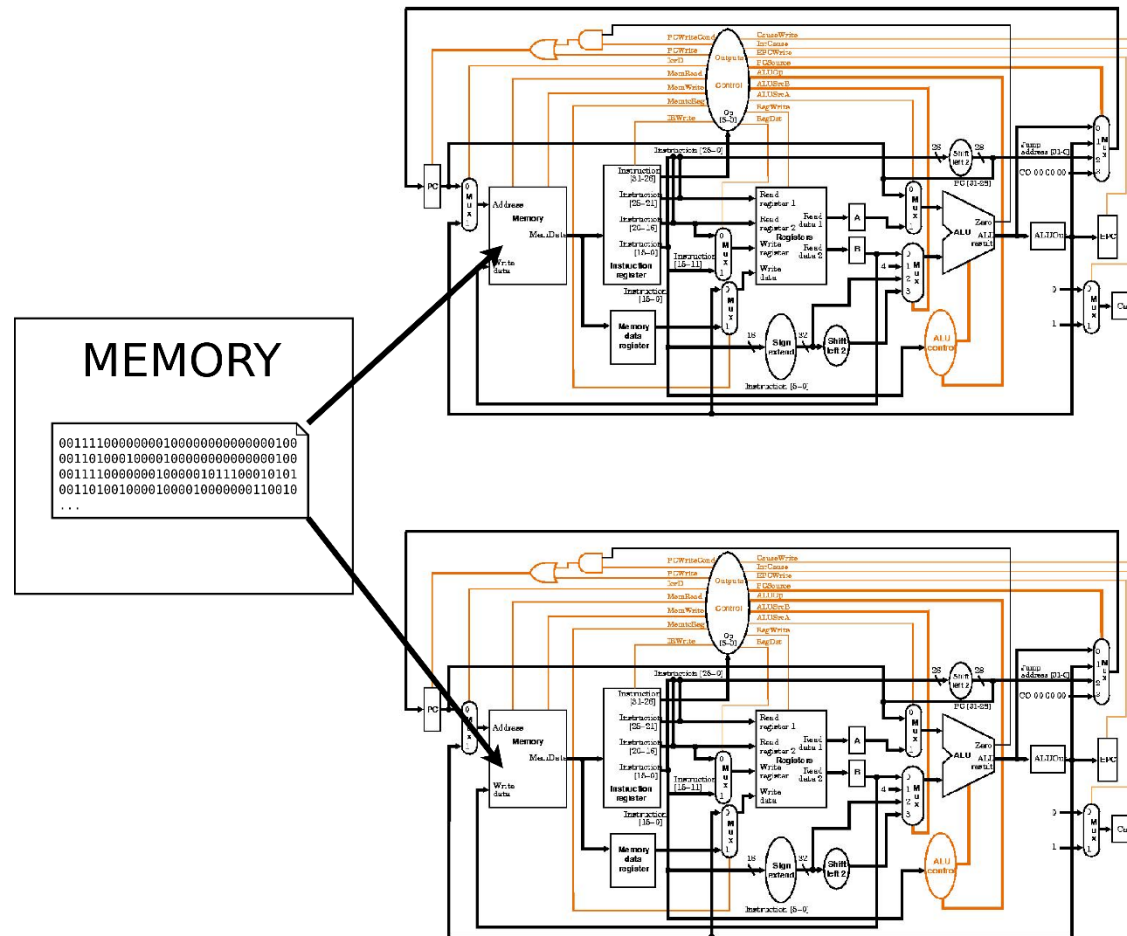
Different types of computer architectures

- Shared memory: all processors access the same memory (e.g. typical server)
 - General model: non-uniform memory access for every processor with protocol for cache coherency (ccNUMA)




Different types of computer architectures

- Shared memory: all processors access the same memory
 - The programmer is seeing a single unified memory
 - Different sequences of execution (threads) run on the same process
 - Different memory modules may be used, but only one logical memory space is addressed
 - Communication between processors is done implicitly



Many threads in a process: OpenMP

- A process creates lightweight instances of itself (threads) that are coordinated for simultaneous execution
 - A thread shares the program code and data section with all other threads inside the same process
 - The parallel computing approach is called fork-join
 - A sequential program begins with one process (that is, only one thread)
 - A parallel region is defined by creating (spawning) threads from the process and destroying the original thread remains



... and this is where we are going to start!

(Break until 10:45)

INTRODUCTION TO PARALLEL PROGRAMMING

Caspar van Leewen
Carlos Teijeiro Barjas